

COMP1001

Computer Systems

Dr. Vasilios Kelefouras

Email: v.kelefouras@plymouth.ac.uk

Website:

<https://www.plymouth.ac.uk/staff/vasilios-kelefouras>

Outline

2

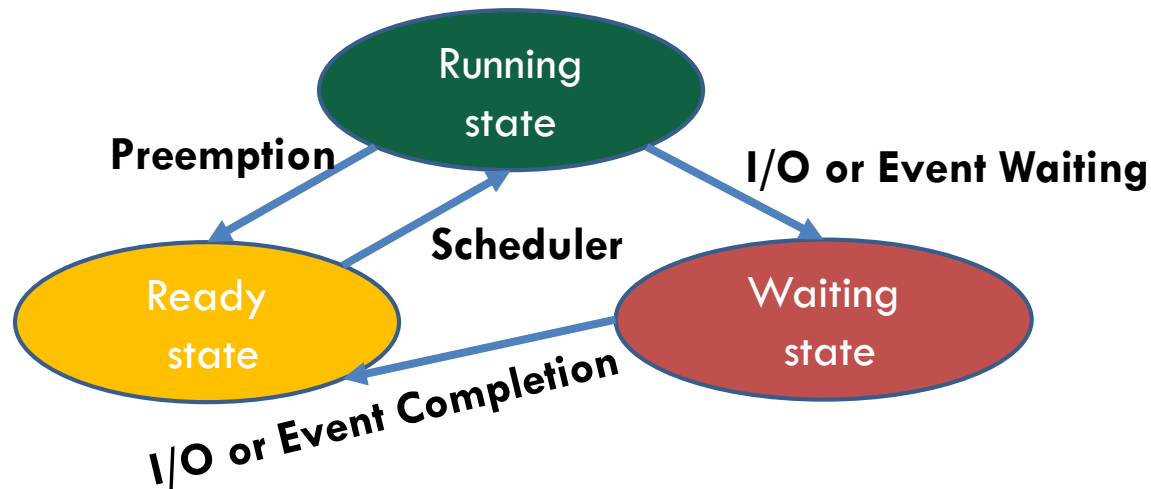
- Processes States
- Context Switch
- System Calls
- Processes
- Cloning a process
- Threads
- Pthreads

Process States (1)

Simplistic point of view

3

- At any point in time, a **process** can be in one of **several states**:
 - ▣ **Running State**: The process is running using CPU and memory resources
 - ▣ **Ready State**: A process is ready to run, but it is not running
 - ▣ **Waiting State**: The process cannot run because it is waiting for some event to occur (or data)

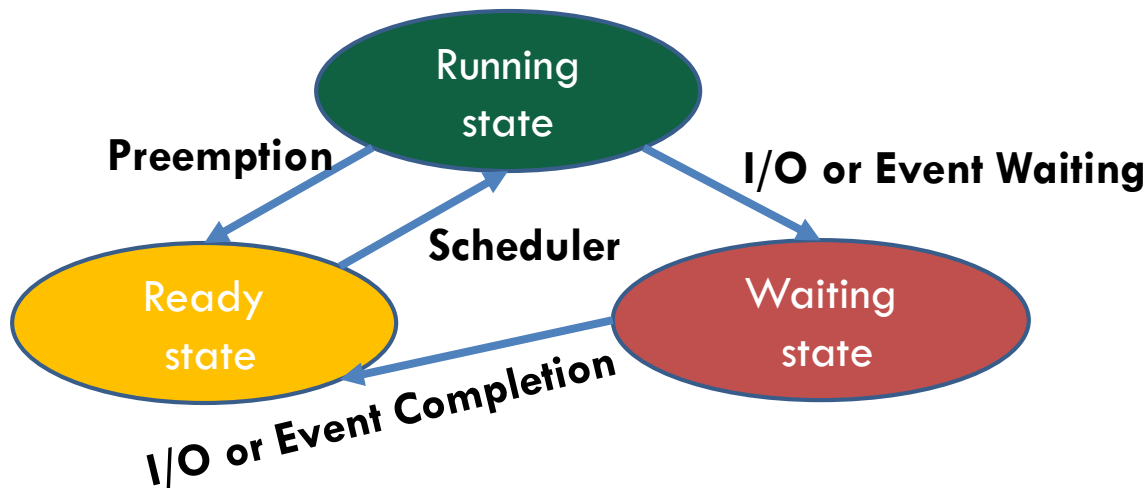


Process States (2)

Simplistic point of view

4

- On a single core CPU, only one process can run at a time, but many processes may be ready for execution or may be blocked
- The OS has a list of the ready and another for the blocked processes
- The scheduler is responsible for choosing the next process to run
- If a running process exceeds a running time limit, then the OS will stop (**pre-empt**) the process and put it back in the ready list. Then, another process will be put to the running state
 - This is called **preemptive multitasking**



Context Switch

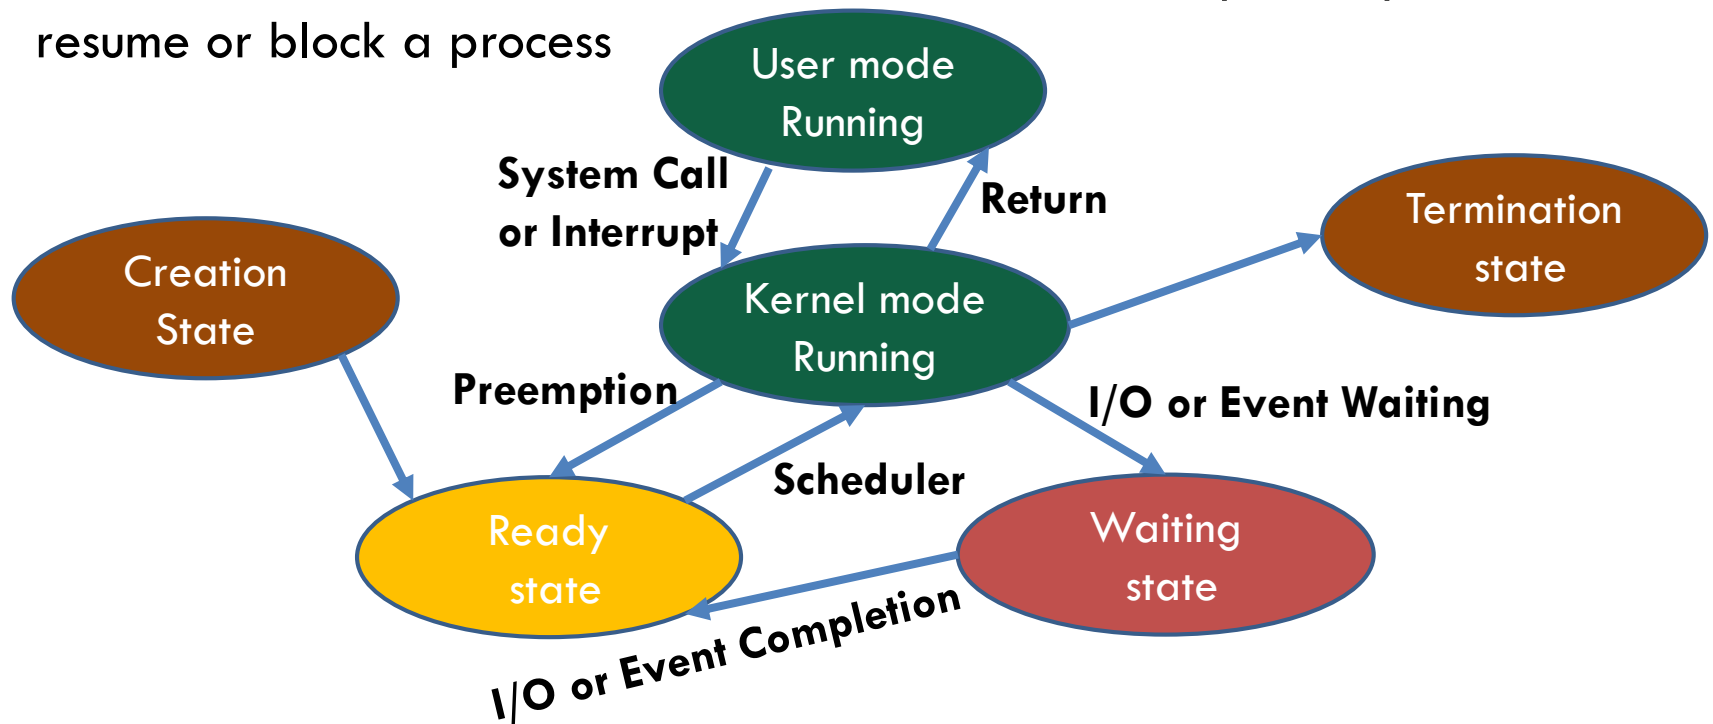
5

- The **scheduler** is responsible for deciding **which process to run** and for **saving and restoring the state of a process** as it gets stopped and when it runs again
- The context of a process is its state
 - ▣ **Text segment** : the compiled code
 - ▣ **Data and bss segments**: all global variables and data structures
 - ▣ **Heap segment** : dynamically allocated data
 - ▣ **Stack segment** : local variables in the running function
 - ▣ **Machine registers**, e.g., Instruction Register
- When a process is pre-empted, the OS does a **context switch**, causing the system to execute in another process context (*there is a performance overhead*)
 - ▣ In a context switch, the OS has **to save all the information needed** so that it can switch back to the earlier process and resume
 - ▣ In a context switch, the OS has to **resume back** the new running process

Process States

6

- **User mode Running** : this is the normal execution of a process. A mode switch to the kernel mode is needed when a system call is performed, e.g., `printf()` or when the process has used its allowed time slot (interrupt) or when an interrupt occurs
- **Kernel mode Running**: it can decide whether to run, pre-empt, terminate, resume or block a process



Scheduling Algorithms

7

- A process scheduler schedules processes using a scheduling algorithm
 - ▣ **First Come, First Served** algorithm
 - ▣ **Shortest Job First** algorithm
 - ▣ **Priority Scheduling** algorithms
 - Each process is assigned a priority
 - Processes with high priority are executed first
 - ▣ **Shortest remaining time** algorithms
 - ▣ **Round Robin** algorithm
 - All processes are assigned to equal time slots and executed one after another
- The aforementioned algorithms can be either
 - ▣ **Pre-emptive**: A process might be pre-empted by a higher priority process
 - ▣ **Non Pre-emptive**: Once a process starts, it cannot stop, until it completes its allotted time

System Calls

8

- ❑ A System Call is a mechanism that provides the interface between a process and the OS.
- ❑ System calls allow user-level processes to request services from the kernel of the OS
- ❑ Programs interact with the OS via system calls
- ❑ Example of system calls are
 - ▣ `fork()`, `exit()`, `open()`, `read()`, `write()`, `sleep()`, `getpid()`, etc
- ❑ System Call steps
 - ▣ The user process runs at user mode until a system call is found
 - ▣ The system call is executed at kernel mode
 - ▣ Control returns to the user process

Note

9

- ❑ The following functions are hard to set up and this is why ***the source code will always be given to you, so don't worry.***
- ❑ The aim of this session is to understand how processes work, not to develop such code

Getting the Process' ID

10

- Every process on the system has a unique process ID number, aka *pid*
 - ▣ *pid* is an integer
 - ▣ You can get pid via the *getpid* system call

```
#include <stdlib.h> //needed for exit()
#include <unistd.h> //needed for getpid()
#include <stdio.h> //needed for printf()
```

```
int main() {
    printf ("my process ID is %d ",getpid() );
    exit(0); //this function terminates the process immediately.}
```

Cloning a process

11

- You can create a process via the **fork system call** which clones a process into two processes running the same code
- **Fork** returns '0' to the child process and the child's '**pid**' to the parent process.
 - ▣ In a failure, -1 is returned.
- No data are shared after **fork()**.
- The process creator is called the parent and the new process is the child
- The parent defines the resources and privileges to its children
- A parent can either wait for the child process to complete or continue in parallel

Cloning a process – 1st Activity

12

```
#include <stdlib.h>    /* needed to define exit() */
#include <unistd.h>    /* needed for fork() and getpid() */
#include <stdio.h>    /* needed for printf() */
```

```
int main( ) {
    int pid;    /* process ID */
```

```
    pid = fork();
    printf("\nHello there\n");
```

```
    if (pid==0) { /* if you are the child process do (a fork returns 0 to the child) */
        printf("I am the child process: pid=%d\n", getpid());    }
```

```
    else if (pid==-1){ /* something went wrong */
        perror("fork"); //this will print an error message
        exit(1);} //terminate the process (1 means not successfully)
```

```
    else{ /* if you are the parent process do (a fork returns a pid to the parent) */
        printf("I am the parent process: pid=%d, child pid=%d\n", getpid(), pid);    }
    exit(0); } //successfully terminate the process
```

Question: What does this program print?

Hello there

I am the parent process: pid11595, child pid=11596

Hello there

I am the child process: pid 11596

2nd Fork activity

What does this program do?

13

```
int main(int argc, char *argv[]) {
    int pid;    /* process ID */
    int a=1, b=0;

    pid = fork();

    a++;
    printf("\nJust after fork a=%d - b=%d\n",a,b);

    if (pid==0) { //if you are the child process
        /* a fork returns 0 to the child */
        printf("I am the child process: pid=%d\n",
getpid());
        a++; b--;
        printf("\nIn child process a=%d\n",a);
        break;
    }

    Else if (pid==-1){ /* something went wrong
        */
        perror("fork");
        exit(1);
    }
    Else { // a fork returns a pid to the parent
        printf("I am the parent process: pid=%d,
child pid=%d\n", getpid(), pid);
        a++; b++;
        printf("\nIn parent process a=%d\n",a);
        break;
    }

    printf("\nJust before end a=%d -
b=%d\n",a,b);

    exit(0);
}
```

2nd Fork activity

What does this program do? (2)

14

//Child Process

int a=1, b=0

a++; //a=2

Just after fork a=2 - b=0

I am the Child process : pid=...

a++; //a=3

b--; //b=-1

Child process a=3

Just before end a=3 - b=-1

//Parent Process

int a=1, b=0

a++; //a=2

Just after fork a=2 - b=0

I am the Parent process : pid=...

a++; //a=3

b++; //b=1

Child process a=3

Just before end a=3 - b=1

execve() system call

15

- ❑ Fork clones the original process
- ❑ Execve() system call replaces the current process with a new one.

//program1.c

```
#include <stdio.h>
```

```
int main(){
```

```
printf("\nHi from Program #1\n");
```

```
}
```

```
gcc program1.c -o program1
```

```
gcc program2.c -o program2
```

```
./program2
```

```
Hi from Program #1
```

//program2.c

```
#include <stdio.h>
```

```
#include <unistd.h> //for execve
```

```
int main(){
```

```
char *args[] = {0};
```

```
execve("program1",args,NULL);
```

```
printf("\n Error\n");
```

```
}
```

Execvp system call

16

- In addition to `execve()` system call, there are other similar routines such as `execvp` and `execvp`.
- `Execvp` allows us to specify all the arguments as parameters to the function and therefore it is easier to use

Program's output:

About to
Hi There

```
int main(int argc, char *argv[]) {
```

```
    char *temp1,*temp2; //these are the memory addresses that the strings are stored in memory
    temp1="Hi";    temp2="There";
```

```
    printf("About to \n");
```

```
    execvp("echo", "e", temp1, temp2, NULL); //the 1st argument is the command to run, the
2nd is the command's name, the others are the arguments passed to execvp. the last must be
NULL
```

```
    perror("execvp");    /* if we get here, execvp failed */
    exit(1);
```

```
}
```


Activity – What does this program do?

17

```
int main() {
    int pid;    /* process ID */

    if (pid = fork()) { //if you are the child process
        funct1 ();
    }

    elseif (pid == -1){ //if you are the parent process
        perror("fork");
        exit(1); }

    else { // a fork returns a pid to the parent
        sleep(5);    /* sleep for 5 seconds */
        printf("Parent is still here!\n");
        break;    }
    exit(0);
}

void funct1 ( ) {

    execlp("echo", "echo_name", "Cheers", "from
child !", (char*)0);
    perror("execlp"); // execlp failed
    exit(1);
}
```

What is difference between thread, process and program?

18

□ **Program:**


- Program is an executable file containing the set of instructions written to perform a specific job on your computer
- For example, *skype.exe* is an executable file containing the set of instructions which help us to run *skype*

□ **Process:**

- Process is an executing instance of a program
- For example, when you double click on the *skype.exe* on your computer, a process is started that will run the skype program

□ **Thread:**

- Thread is the smallest executable unit of a process
- For example, when you run skype program, OS creates a process and starts the execution of the main thread of that process
- A process can have multiple threads
- All threads of the same process share memory of that process



Further Reading

19

- ❑ Chapter 3 and chapter 4 in Operating Systems, Internals and Design Principles, available at https://dinus.ac.id/repository/docs/ajar/Operating_System.pdf
- ❑ POSIX Threads Programming, available at <https://computing.llnl.gov/tutorials/pthreads/>
- ❑ POSIX thread (pthread) libraries available at <https://www.cs.cmu.edu/afs/cs/academic/class/15492-f07/www/pthreads.html>

Thank you