

COMP1001

Computer Systems

Dr. Vasilios Kelefouras

Email: v.kelefouras@plymouth.ac.uk

Website:

<https://www.plymouth.ac.uk/staff/vasilios-kelefouras>

Outline

2

- Introduction to Threads
- Two simple multi-threaded examples

What is difference between thread, process and program?

3

□ Program:


- Program is an executable file containing the set of instructions written to perform a specific job on your computer
- For example, *skype.exe* is an executable file containing the set of instructions which help us to run *skype*

□ Process:

- Process is an executing instance of a program
- For example, when you double click on the *skype.exe* on your computer, a process is started that will run the *skype* program

□ Thread:

- Thread is the smallest executable unit of a process
- For example, when you run *skype* program, OS creates a process and starts the execution of the main thread of that process
- A process can have multiple threads
- All threads of the same process share memory of that process



Processes and Threads

4

- A process is a program in memory along with its dynamically allocated storage (heap), its stack storage and its execution context (state of registers and instruction pointer)
- A process might have more than one threads – multithreaded process
- A process can be broken down into
 1. Text segment, heap segment, static data segments
 2. Stack , instruction pointer and registers
 - Each thread has its own stack, instruction pointer and registers
 - Each thread has its own thread ID

Threads – Advantages of using threads

5



- Threads are light compared to processes
 - ▣ The OS does not have to create a new memory map (recall previous lecture) for a new thread, as it does for a process
 - ▣ The OS does not have to keep track of open files amongst different threads
 - ▣ Switching between threads of the same process has less overhead than process switch
 - ▣ A multithreading application scales better by adding more CPU cores
- Programming is easier as all the threads share global variables

How can we parallelize this program multiple threads?

See sqrt.c program on github

6

```
#define N 1000  
float X[N], Y[N];  
int i, j;  
  
for (j=0; j<N; j++)  
    Y[i] = sqrt( X[j] );
```



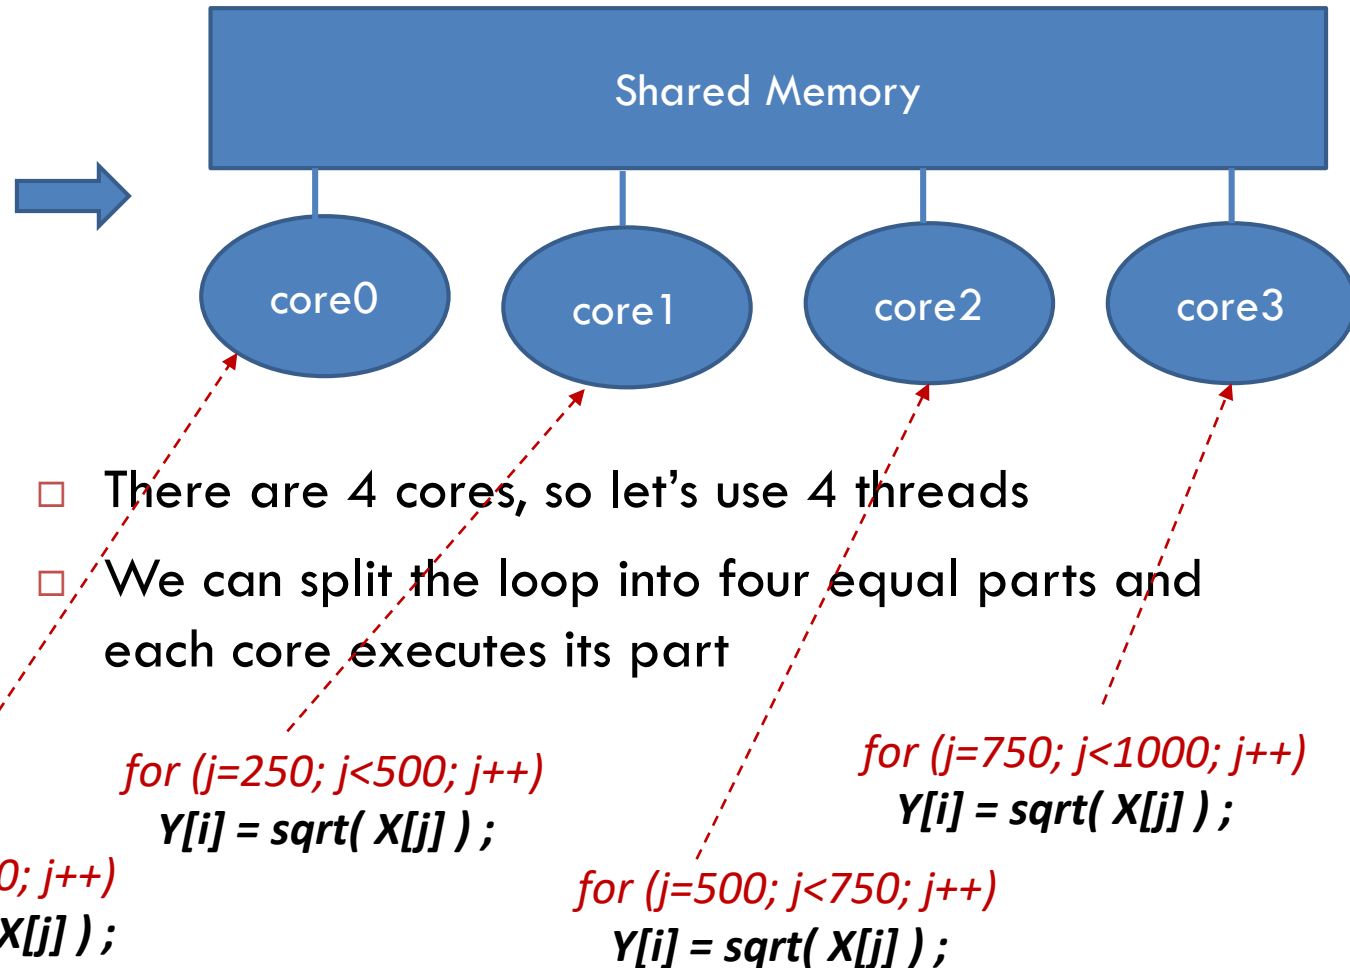

How can we parallelize this program multiple threads?

See sqrt.c program on github

7

```
#define N 1000  
float X[N], Y[N];  
int i, j;
```

```
for (j=0; j<N; j++)  
    Y[i] = sqrt( X[j] );
```

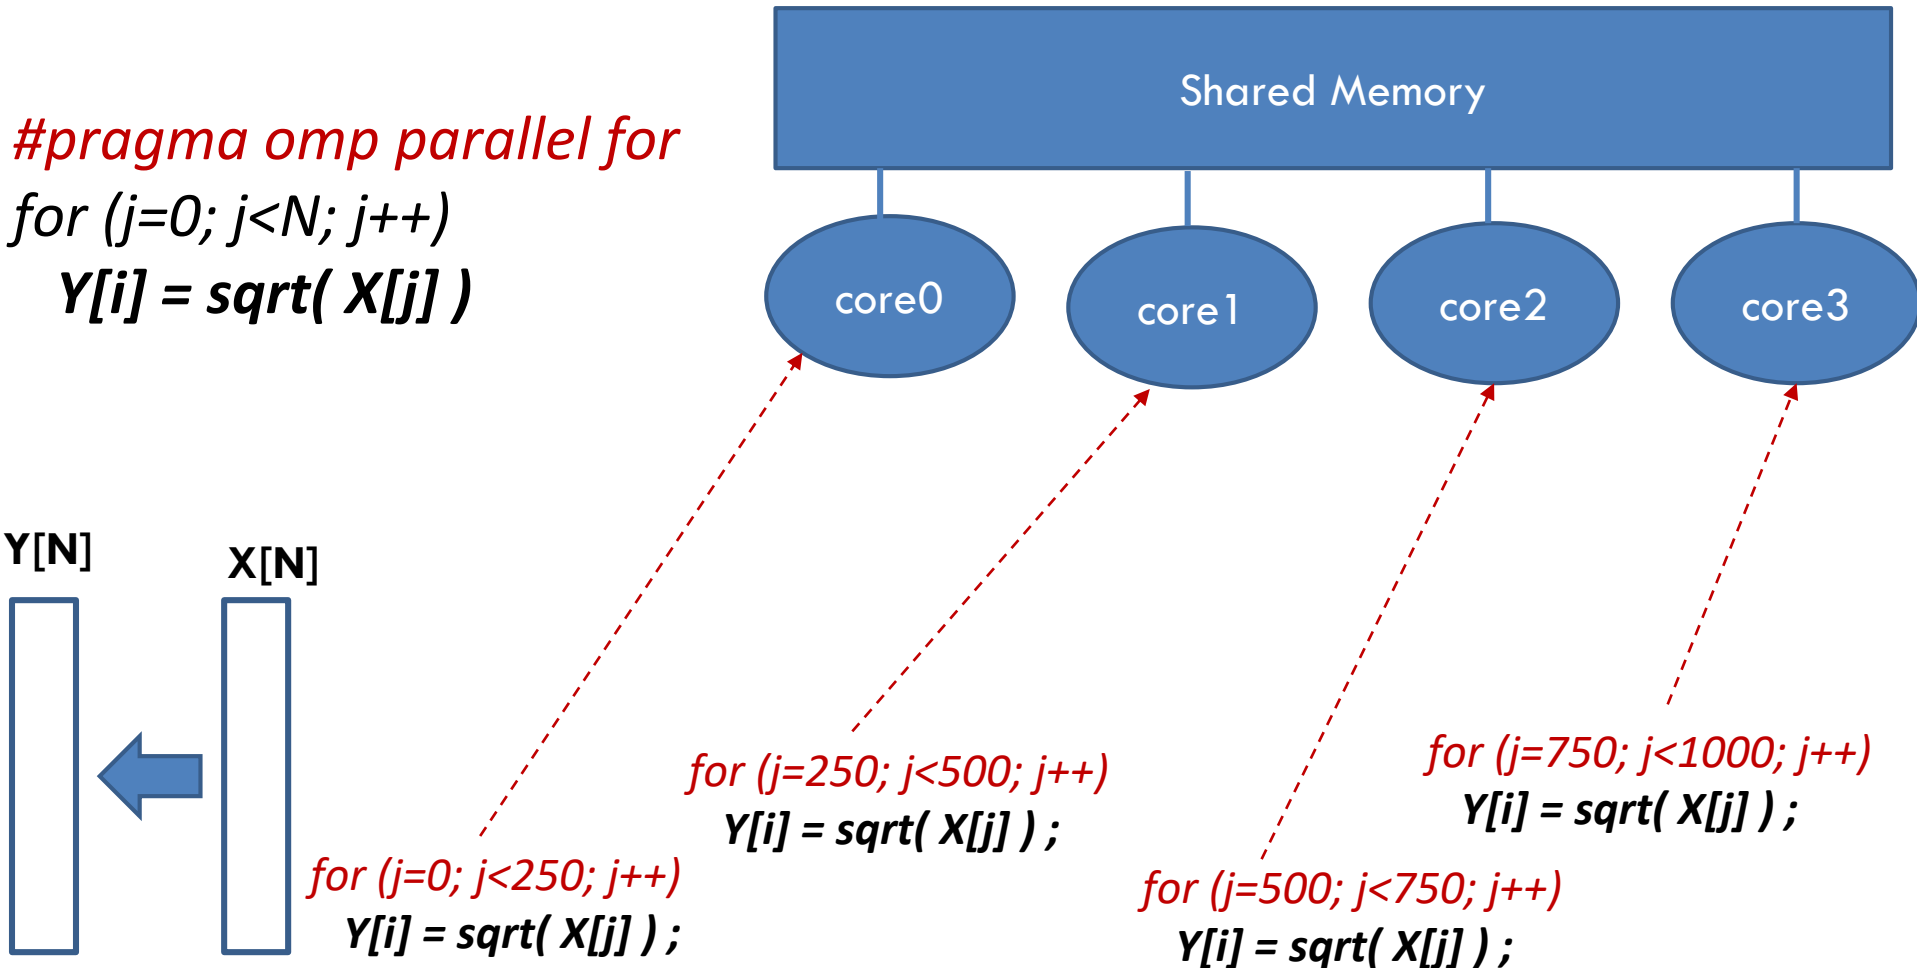


How can we parallelize this program multiple threads?

See sqrt.c program on github

8

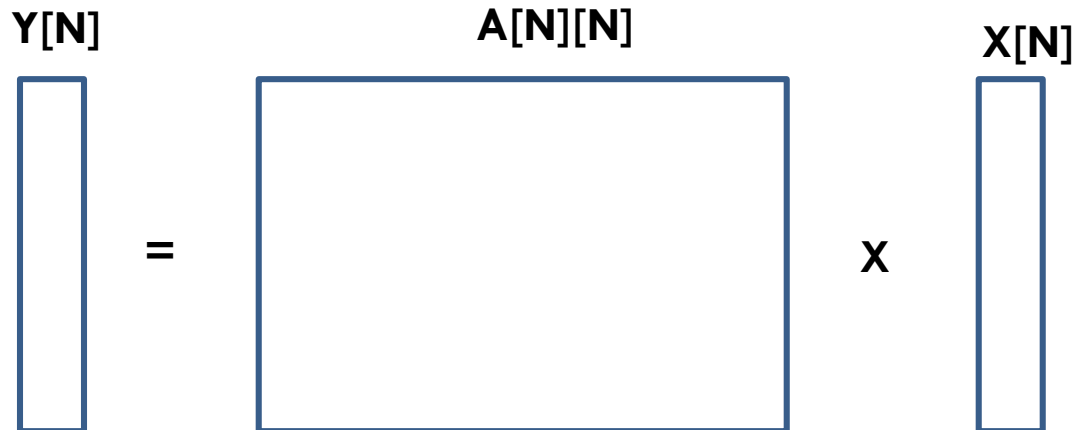

- This program yields the same behaviour as the previous one



How can we parallelize this program?


9

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



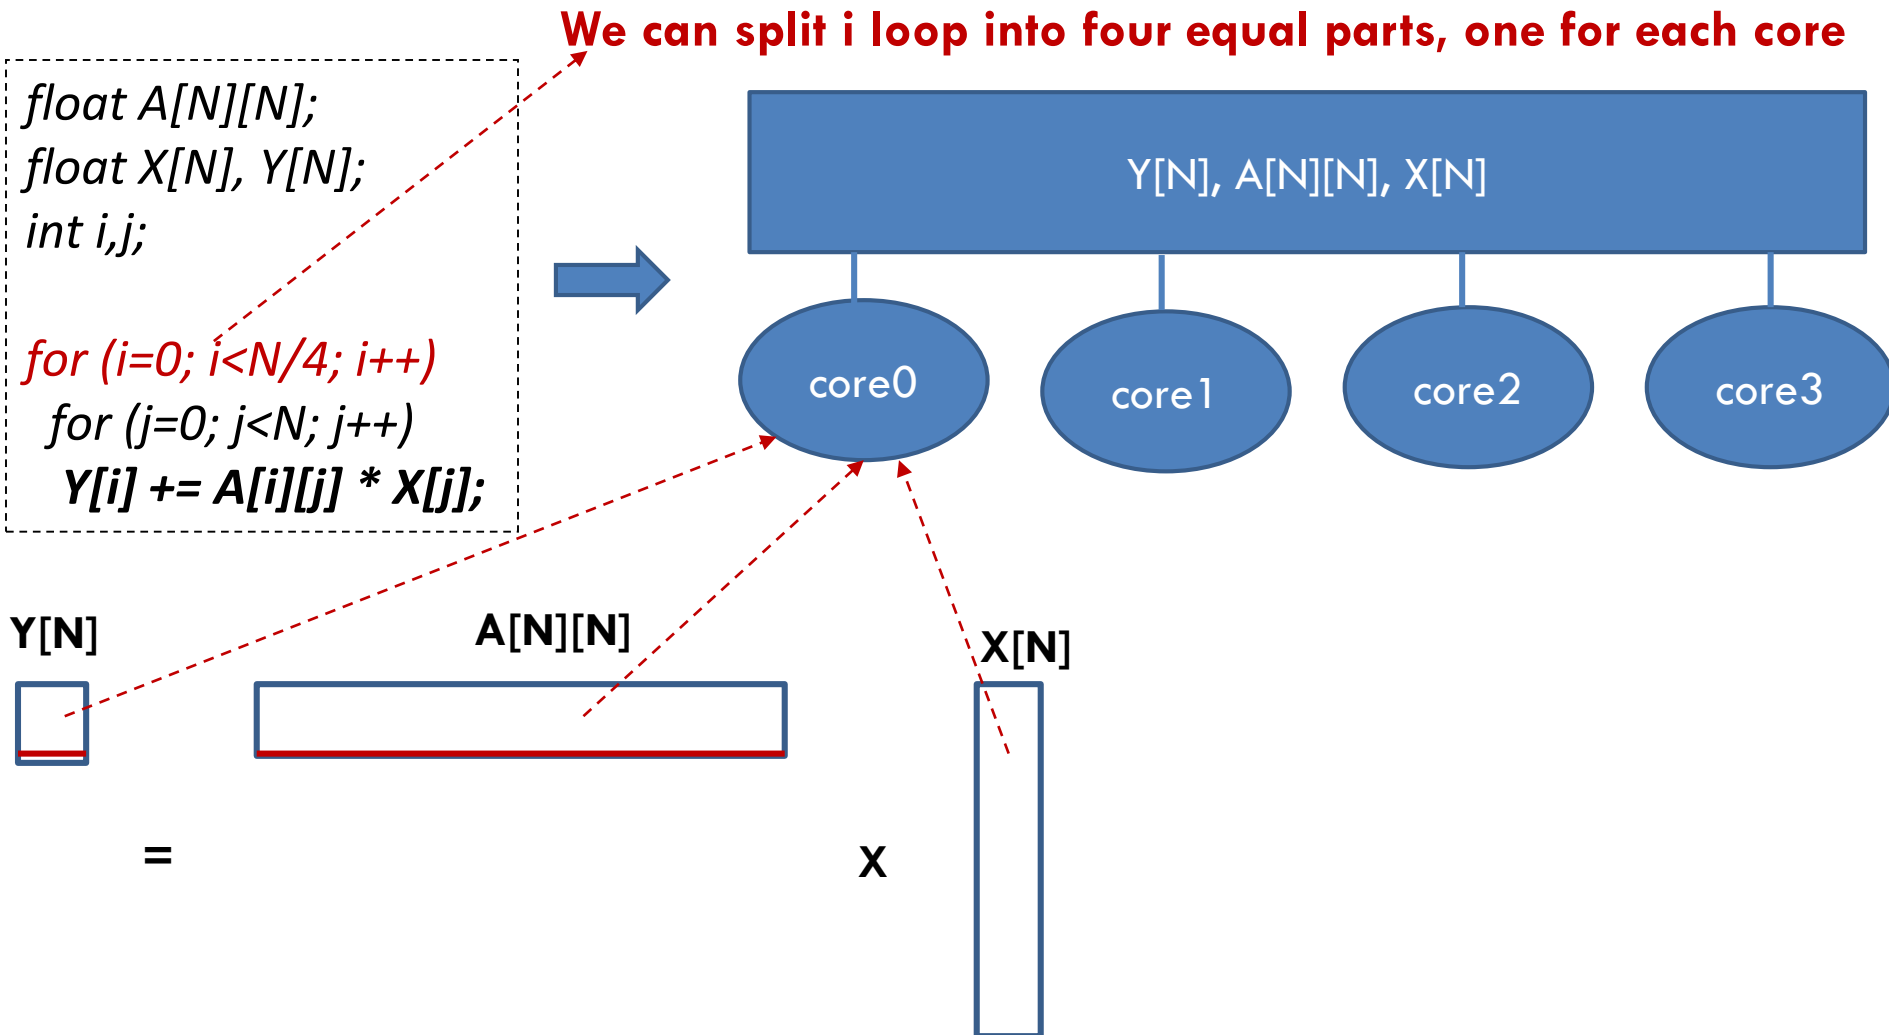
How can we parallelize this program?

10



How can we parallelize this program?

11



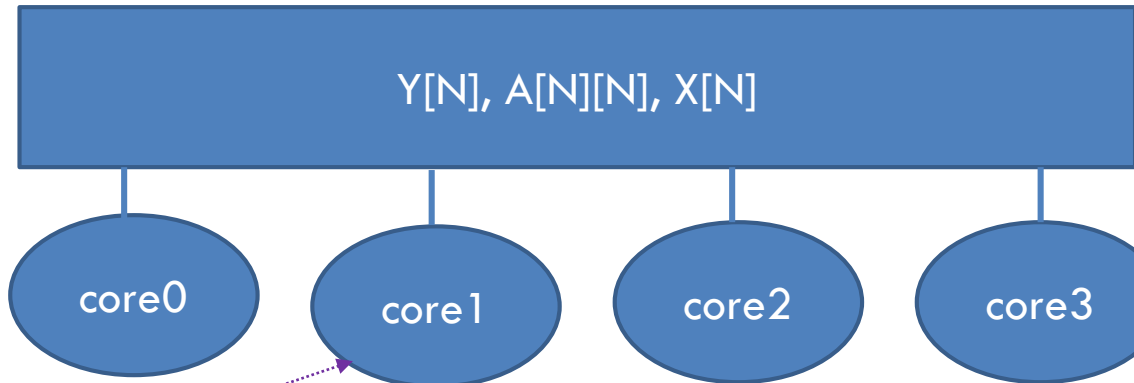
How can we parallelize this program?

12

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=N/4; i<N/2; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```


We can split i loop into four equal parts, one for each core



$Y[N]$

$A[N][N]$

$X[N]$



=



X

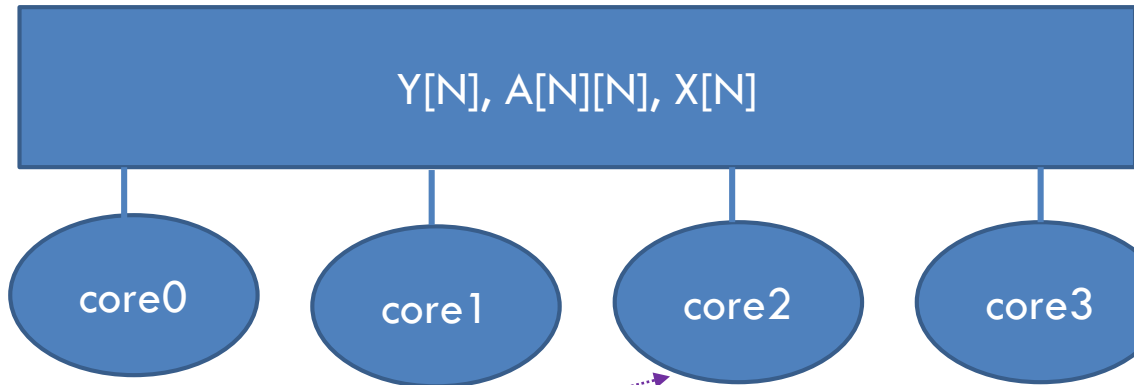
How can we parallelize this program?

13

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=N/2; i<3*N/4; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```

We can split i loop into four equal parts, one for each core



$Y[N]$

$A[N][N]$

$X[N]$



=



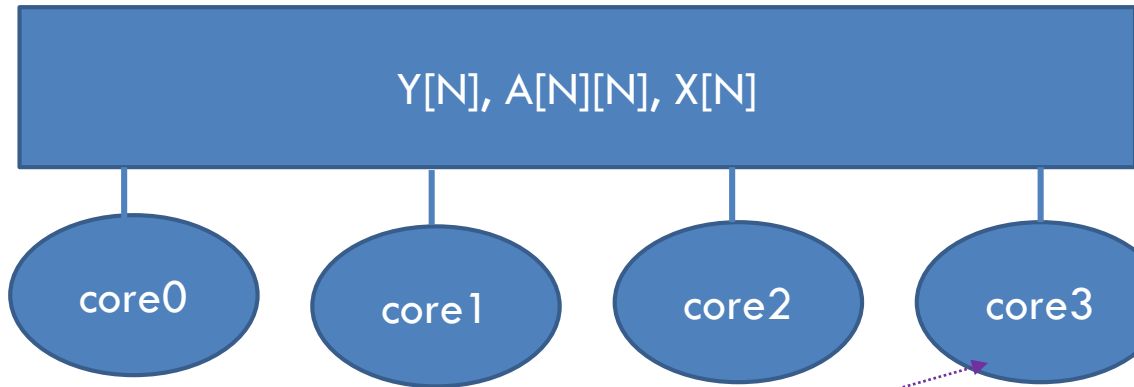
X

How can we parallelize this program?

14

We can split i loop into four equal parts, one for each core

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=3*N/4; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



$Y[N]$

$A[N][N]$

$X[N]$

=


X



How can we parallelize this program?

15

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
#pragma omp parallel  
for private (j)  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



Performance speedup

	N=100	N=200	N=500	N=1000	N=2000
2 threads	x1.47	x1.97	x1.98	x1.99	x1.99
4 threads	x1.16	x2.3	x3.47	x3.69	x3.7

Why our code does not scale well in these cases?

How can we parallelize this program?

16

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
#pragma omp parallel  
for private (j)  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```

- When using OpenMP, there is **an overhead in creating and synchronizing the threads**
- When this overhead becomes comparable to the thread's computation, the speedup is low
 - ▣ This is why the code does not scale well for small input sizes

Performance speedup

	N=100	N=200	N=500	N=1000	N=2000
2 threads	x1.47	x1.97	x1.98	x1.99	x1.99
4 threads	x1.16	x2.3	x3.47	x3.69	x3.700

Why our code does not scale well in these cases?

Further Reading

17

- Chapter 3 and chapter 4 in Operating Systems, Internals and Design Principles, available at https://dinus.ac.id/repository/docs/ajar/Operating_System.pdf
- POSIX Threads Programming, available at <https://computing.llnl.gov/tutorials/pthreads/>
- POSIX thread (pthread) libraries available at <https://www.cs.cmu.edu/afs/cs/academic/class/15492-f07/www/pthreads.html>

Thank you