

# COMP1001

## Computer Systems

Dr. Vasilios Kelefouras

Email: [v.kelefouras@plymouth.ac.uk](mailto:v.kelefouras@plymouth.ac.uk)

Website:

<https://www.plymouth.ac.uk/staff/vasilios-kelefouras>

# Outline

2

- Introduction to Threads
- Introduction to Pthreads
- Introduction to OpenMP

# What is difference between thread, process and program?

3

## □ Program:


- Program is an executable file containing the set of instructions written to perform a specific job on your computer
- For example, *skype.exe* is an executable file containing the set of instructions which help us to run *skype*

## □ Process:

- Process is an executing instance of a program
- For example, when you double click on the *skype.exe* on your computer, a process is started that will run the *skype* program

## □ Thread:

- Thread is the smallest executable unit of a process
- For example, when you run *skype* program, OS creates a process and starts the execution of the main thread of that process
- A process can have multiple threads
- All threads of the same process share memory of that process



# Processes and Threads

4

- A process is a program in memory along with its dynamically allocated storage (heap), its stack storage and its execution context (state of registers and instruction pointer)
- A process might have more than one threads – multithreaded process
- A process can be broken down into
  1. Text segment, heap segment, static data segments
  2. Stack , instruction pointer and registers
    - Each thread has its own stack, instruction pointer and registers
    - Each thread has its own thread ID

# Threads – Advantages of using threads

5

- Threads are light compared to processes
  - ▣ The OS does not have to create a new memory map (recall previous lecture) for a new thread, as it does for a process
  - ▣ The OS does not have to keep track of open files amongst different threads
  - ▣ Switching between threads of the same process has less overhead than process switch
  - ▣ A multithreading application scales better by adding more CPU cores
- Programming is easier as all the threads share global variables

# POSIX Threads

6

- ❑ The POSIX thread libraries are a standards based thread API for C/C++.
- ❑ It allows to create a multiprocessing flow.
- ❑ It is most effective on multi-core systems where the process flow can be scheduled to run on multiple cores and thus improve performance.
- ❑ Creating multiple threads require less overhead than "forking" or spawning a new process because the system does not create a new memory map.
- ❑ Performance gains are also found on uniprocessor systems which exploit latency in I/O and other system functions which may halt process execution.
  - ▣ One thread may execute while another is waiting for I/O or some other system latency

# Create a new Thread using pthreads

7

- Pthreads are **defined** by using:

- ▣ **Pthread\_t** my\_thread;
- ▣ We must include the **#include <pthread.h>** library

- A thread is **created** via

- ▣ **Pthread\_t** my\_thread;  
**pthread\_create**(&my\_thread, NULL, func, arg);

- ▣ int **pthread\_create**(pthread\_t \* thread, *//this is the thread ID*  
**const pthread\_attr\_t \* attr**, *//we will set this to NULL (default thread attributes)*  
**void \* (\*start\_routine)(void \*)**, *//pointer to the function to be threaded.*  
*Function has a single argument: pointer to void.*  
**void \*arg**); *//pointer to argument of function. To pass multiple arguments, send a pointer to a structure.*

# Pthreads - Exit and Join

8

- **Exit a thread.** A thread can exit by
  - ▣ calling *pthread\_exit* or
  - ▣ *pthread\_join()*
  
- **Join two threads**
  - ▣ By calling Pthread\_join( ) command, the program will wait here until the thread finishes

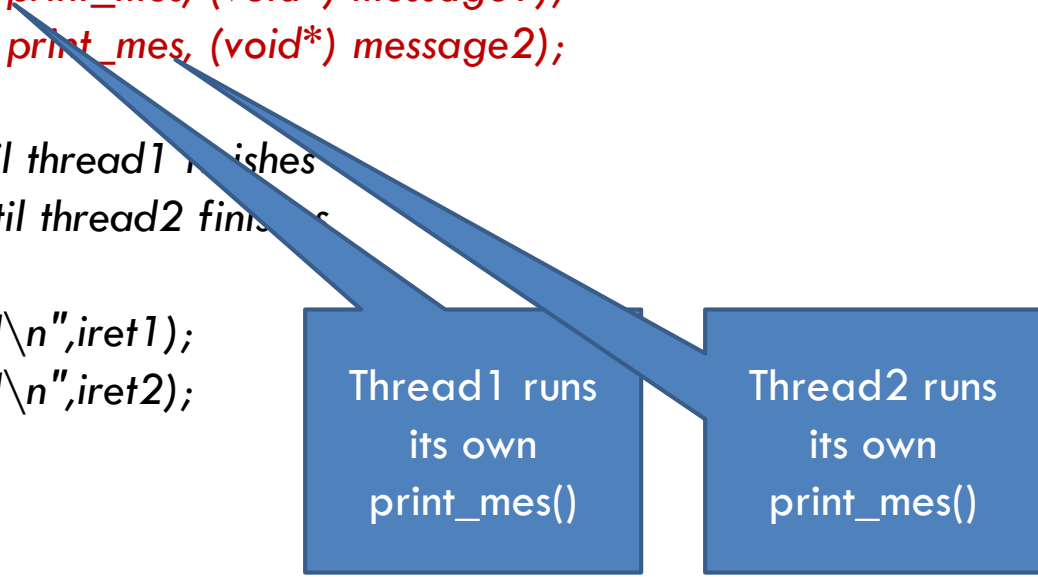


# Pthreads example (1)

9

```
int main() {  
  
    pthread_t thread1, thread2, thread3, thread4;  
  
    char *message1 = "Thread 1";  
    char *message2 = "Thread 2";  
    int iret1, iret2;  
  
    iret1 = pthread_create( &thread1, NULL, print_mes, (void*) message1);  
    iret2 = pthread_create( &thread2, NULL, print_mes, (void*) message2);  
  
    pthread_join( thread1, NULL); //wait until thread1 finishes  
    pthread_join( thread2, NULL); //wait until thread2 finishes  
  
    printf("Thread 1 completed, returned: %d\n",iret1);  
    printf("Thread 2 completed, returned: %d\n",iret2);  
  
    return 0; }
```

```
void *print_mes( void *ptr )  
{  
    char *message;  
    message = (char *) ptr;  
    printf("%s \n", message);  
}
```



Thread1 runs  
its own  
print\_mes()

Thread2 runs  
its own  
print\_mes()

# Pthreads example (2)

10

```
void *print_mes ( void *ptr )  
{  
    char *message;  
    message = (char *) ptr;  
    printf("%s \n", message);  
}
```

- ❑ this function takes as input a pointer to void (pointer to anything).
- ❑ It returns a pointer to void too.
- ❑ A void pointer is a pointer that has no associated data type with it.
- ❑ A void pointer can hold address of any type and can be typecasted to any type.

# Pthreads example (3)

## What does this program print?

11

- The 1<sup>st</sup> time we run the program it generates:
- The 2<sup>nd</sup> time we run the program it generates:
- *Each thread runs its own copy of the function, **at any order.***

Program's output:

**Thread1**

**Thread2**

**Thread1 completed, returned: 0**

**Thread2 completed, returned: 0**

Program's output:

**Thread2**

**Thread1**

**Thread1 completed, returned: 0**

**Thread2 completed, returned: 0**

# What if we delete the join() function?

12

```
int main() {  
  
    pthread_t thread1, thread2, thread3, thread4;  
  
    char *message1 = "Thread 1";  
    char *message2 = "Thread 2";  
    int iret1, iret2;  
  
    iret1 = pthread_create( &thread1, NULL, print_message_function, (void*) message1); //create  
    thread1 which will run its own print_message_function()  
    iret2 = pthread_create( &thread2, NULL, print_message_function, (void*) message2);  
  
    pthread_join( thread1, NULL);  
  
    pthread_join( thread2, NULL);  
  
    printf("Thread 1 completed, returned: %d\n",iret1);  
    printf("Thread 2 completed, returned: %d\n",iret2);  
  
    return 0; }
```

```
void *print_message_function( void *ptr )  
{  
    char *message;  
    message = (char *) ptr;  
    printf("%s \n", message);  
}
```

Program's output:

**Thread1 completed, returned: 0**

**Thread2 completed, returned: 0**

**Thread2**

**Thread1**



**Wrong output**

# How can we parallelize this program using pthreads?

## sqrt.c program

13

```
#define N 1000  
float X[N], Y[N];  
int i, j;  
  
for (j=0; j<N; j++)  
    Y[j] = sqrt( X[j] );
```



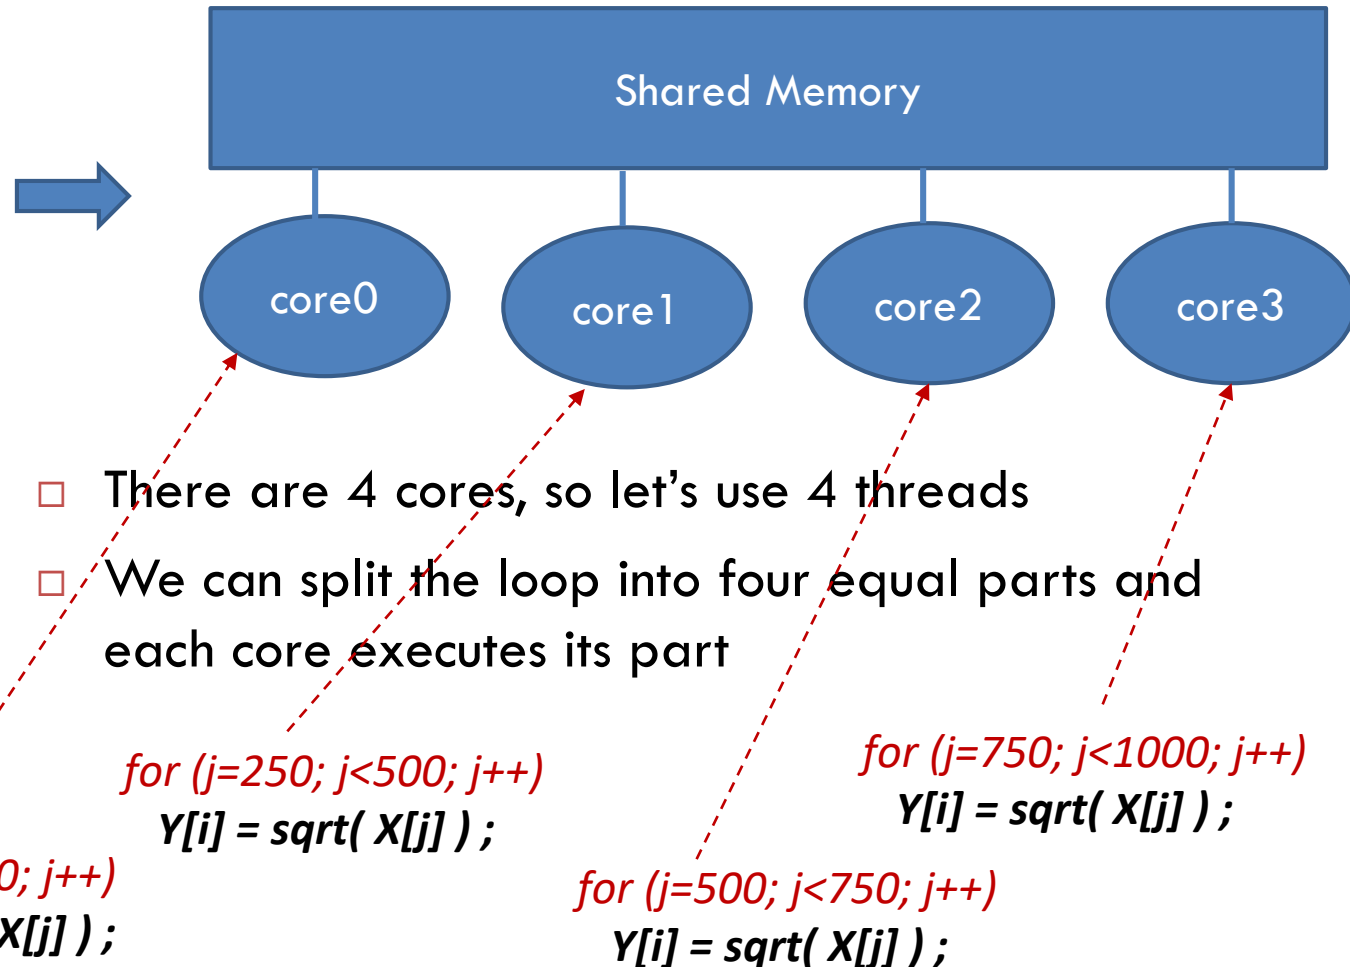

# How can we parallelize this program using pthreads?

## sqrt.c program

14

```
#define N 1000
float X[N], Y[N];
int i, j;

for (j=0; j<N; j++)
    Y[i] = sqrt( X[j] );
```



# How can we parallelize this program using pthreads?

## sqrt.c program

15

```
...  
for (thread_num = 0; thread_num < NUM_THREADS; thread_num++)  
    pthread_create(&thread_handles[thread_num], NULL, sqrt, (void*) thread_num);  
.....
```

- Create '**NUM\_THREADS**' threads where they all execute **sqrt()** function
- Let's assume  $N=1000$  and  $\text{NUM\_THREADS}=4$

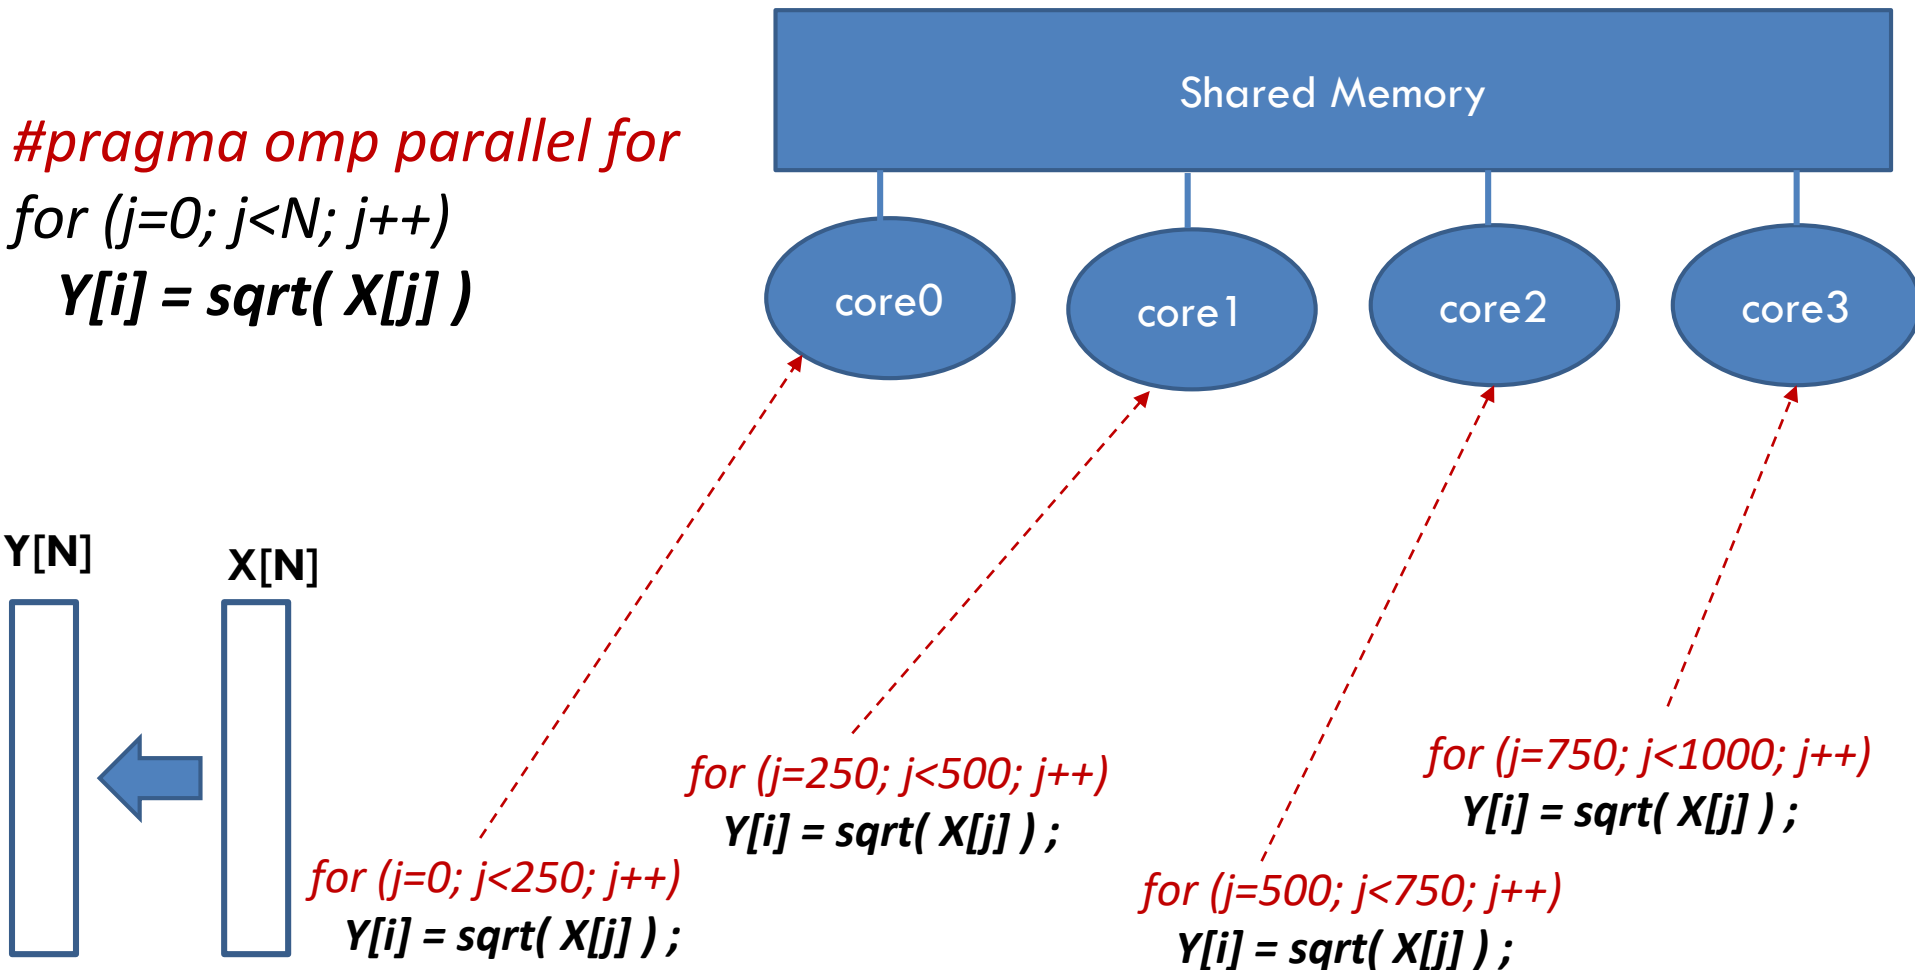
```
void *sqrt (void* thread_num) {  
  
    long my_rank = (long) thread_num;  
    int i, j;  
    int local = N/NUM_THREADS;  
  
    int starting_row = my_rank * local;  
    int ending_row = starting_row + local - 1;  
  
    for (i = starting_row; i <= ending_row; i++) {  
        Y[i] = sqrt ( X[j] );    }  
  
    return 0;  
}
```

Thread	Starting row	Ending row
0	0	249
1	250	499
2	500	749
3	750	999

# How can we parallelize the sqrt program by using OpenMP framework ?

16

- This program yields the same behaviour as the previous one



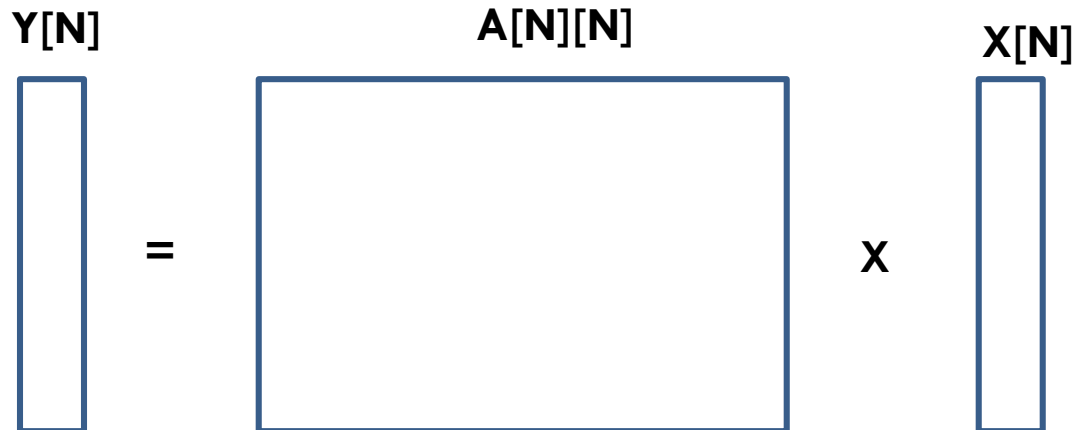



# How can we parallelize this program using pthreads?

## MVM\_pthreads.c program

17


```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



# How can we parallelize this program using pthreads?

## MVM\_pthreads.c program

18



# How can we parallelize this program using pthreads?

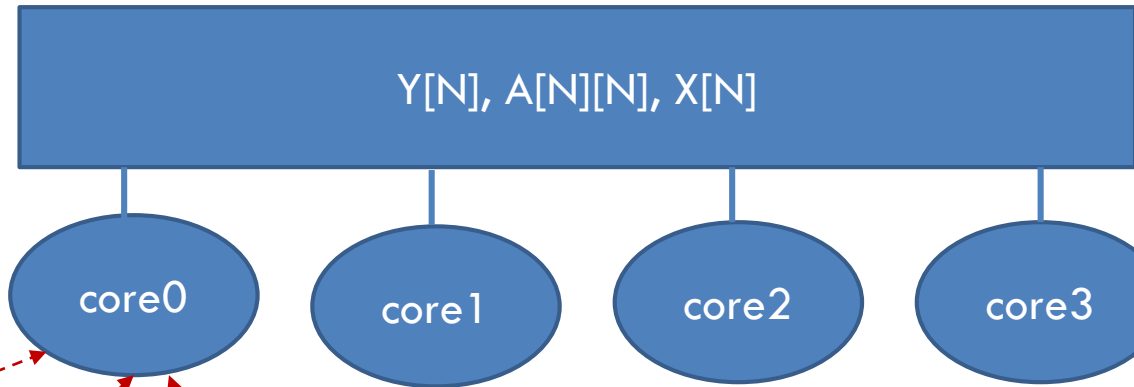
## MVM\_pthreads.c program

19

**We can split i loop into four equal parts, one for each core**

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=0; i<N/4; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



$Y[N]$

$A[N][N]$

$X[N]$



=

X

# How can we parallelize this program using pthreads?

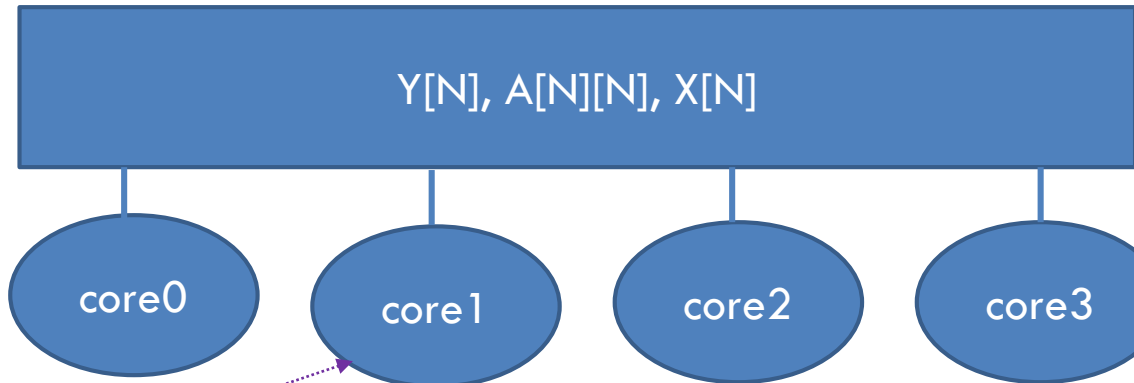
## MVM\_pthreads.c program

20

**We can split i loop into four equal parts, one for each core**

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```


```
for (i=N/4; i<N/2; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



`Y[N]`

`A[N][N]`

`X[N]`



=



`X`



# How can we parallelize this program using pthreads?

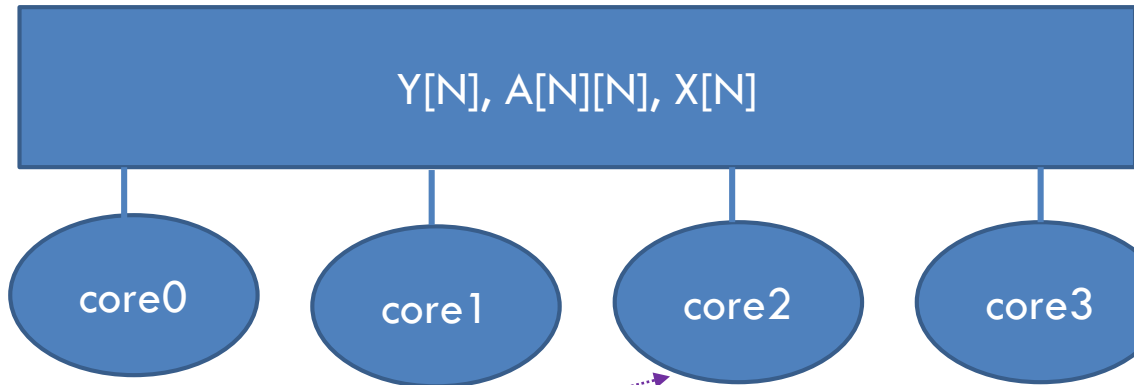
## MVM\_pthreads.c program

21

**We can split i loop into four equal parts, one for each core**

```
float A[N][N];  
float X[N], Y[N];  
int i,j;
```

```
for (i=N/2; i<3*N/4; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



$Y[N]$

$A[N][N]$

$X[N]$



=



$X$

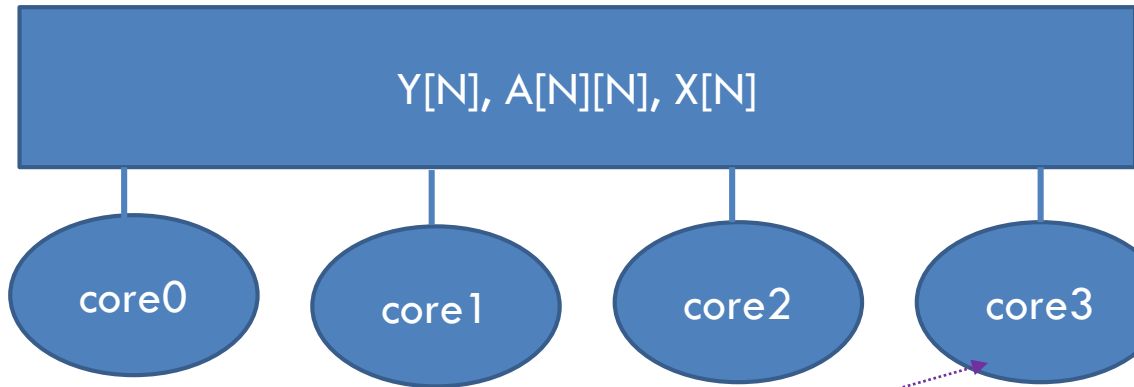
# How can we parallelize this program using pthreads?

## MVM\_pthreads.c program

22

**We can split i loop into four equal parts, one for each core**

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=3*N/4; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



`Y[N]`

`A[N][N]`

`X[N]`

=

X



# How can we parallelize this program using pthreads?

## MVM\_pthreads.c program

23

...

```
for (thread_num = 0; thread_num < NUM_THREADS; thread_num++)  
    pthread_create(&thread_handles[thread_num], NULL, MVM, (void*) thread_num);  
.....
```

.....

```
void *MVM (void* rank) {
```

```
    long my_rank = (long) rank;  
    int i, j;  
    int local = N/NUM_THREADS;
```

```
    int starting_row = my_rank * local;  
    int ending_row = starting_row + local - 1;
```

```
    for (i = starting_row; i <= ending_row; i++) {  
        for (j = 0; j < N; j++)  
            Y[i] += A[i][j] * X[j];    }
```

```
    return 0;
```

```
}
```

- Create 'NUM\_THREADS' threads where they all execute MVM() function
- Let's assume N=1000 and NUM\_THREADS=4


Thread	Starting row	Ending row
0	0	249
1	250	499
2	500	749
3	750	999

# How can we parallelize this program using pthreads?

## Results from the lab session

24

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    Y[i] += A[i][j] * X[j];
```



Performance speedup

	N=100	N=200	N=500	N=1000	N=2000
2 threads	x1.47	x1.97	x1.98	x1.99	x1.99
4 threads	x1.16	x2.3	x3.47	x3.69	x3.7

**Why our code does not scale well in these cases?**



# How can we parallelize this program using pthreads?

## Results from the lab session

25

```
float A[N][N];  
float X[N], Y[N];  
int i,j;  
  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        Y[i] += A[i][j] * X[j];
```

- When using pthreads, there is **an overhead in creating and synchronizing the threads**
- When this overhead becomes comparable to the thread's computation, the speedup is low
  - ▣ This is why the code does not scale well for small input sizes

Performance speedup

	N=100	N=200	N=500	N=1000	N=2000
2 threads	x1.47	x1.97	x1.98	x1.99	x1.99
4 threads	x1.16	x2.3	x3.47	x3.69	x3.700

**Why our code does not scale well in these cases?**

# OpenMP as an easier and higher level solution

26

- Using Pthreads is not that easy
  - ▣ Low level API
- OpenMP provides a higher level API which is easier to use
- See below how our previous example looks like using openMP...
- Pthreads provide more flexibility

```
#pragma omp parallel for private (i,j)  
for (i=0; i<N; i++)  
for (j=0; j<N; j++)  
Y[i] += A[i][j] * X[j];
```

# Further Reading

27

- Chapter 3 and chapter 4 in Operating Systems, Internals and Design Principles, available at [https://dinus.ac.id/repository/docs/ajar/Operating\\_System.pdf](https://dinus.ac.id/repository/docs/ajar/Operating_System.pdf)
- POSIX Threads Programming, available at <https://computing.llnl.gov/tutorials/pthreads/>
- POSIX thread (pthread) libraries available at <https://www.cs.cmu.edu/afs/cs/academic/class/15492-f07/www/pthreads.html>

Thank you